

大数据分析

Scalable Machine Learning
decision tree

刘盛华

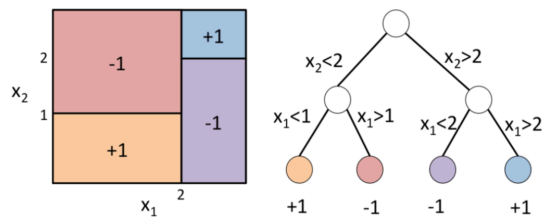
Outline

- Decision Tree
- Random Forest
- Gradient Boosted Decision Tree (GBDT)

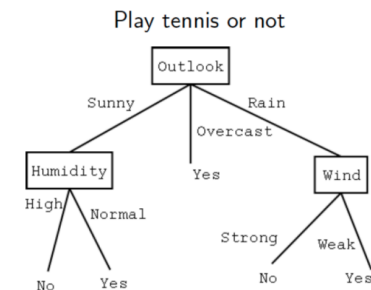
Cho-Jui Hsieh, UC Davis ECS171: Machine Learning

Decision Tree

- Each node checks one feature x_i :
 - Go left if $x_i < \text{threshold}$
 - Go right if $x_i \geq \text{threshold}$



A real example



Decision Tree

Strength:

- it's a **nonlinear** classifier
- Better **interpretability**
- Can naturally handle **categorical** features

非线性
更好的解释性
自然处理分类特征

Computation:

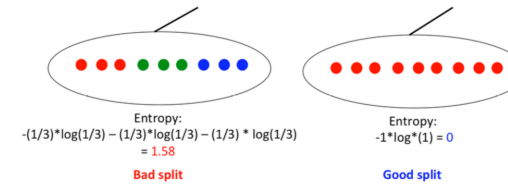
- Training: **slow**
- Prediction: **fast**
- h operations (h : depth of the tree, usually ≤ 15)

Splitting the node

- Classification tree: Split the node to maximize entropy
- Let S be set of data points in a node, $c = 1, \dots, C$ are labels:

$$\text{Entropy} : H(S) = - \sum_{c=1}^C p(c) \log p(c),$$

- where $p(c)$ is the proportion of the data belong to class c .
 - Entropy=0 if all samples are in the same class
 - Entropy is large if $p(1) = \dots = p(C)$



Information Gain

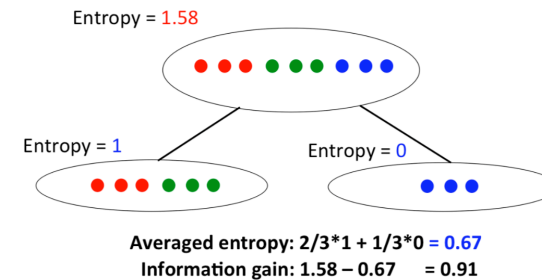
- The averaged entropy of a split $S \rightarrow S_1, S_2$

$$\frac{|S_1|}{|S|} H(S_1) + \frac{|S_2|}{|S|} H(S_2)$$

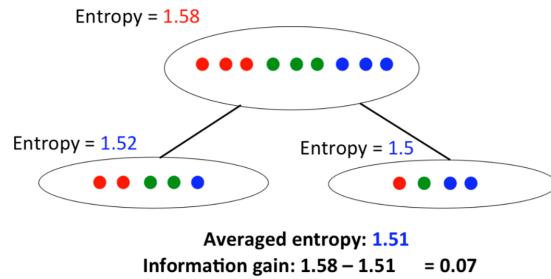
- Information gain: measure how good is the split

$$H(S) - \left(\frac{|S_1|}{|S|} H(S_1) + \frac{|S_2|}{|S|} H(S_2) \right)$$

Information Gain



Information Gain

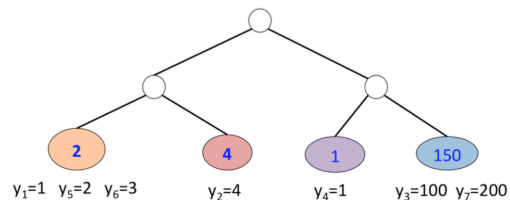


Splitting the node

- Given the current node, how to find the **best split**?
- For all the **features** and all the **threshold**
 - Compute the information gain after the split
 - Choose the best one (**maximal information gain**)
- For n samples and d features: need $O(nd)$ time

Regression Tree

- Assign a real number for each leaf
- Usually **averaged y values** for each leaf (minimize square error)



Regression Tree

Objective function:

$$\min_F \frac{1}{n} \sum_{i=1}^n (y_i - F(\mathbf{x}_i))^2 + (\text{Regularization})$$

The quality of partition $S = S_1 \cup S_2$ can be computed by the objective function:

$$\sum_{i \in S_1} (y_i - y^{(1)})^2 + \sum_{i \in S_2} (y_i - y^{(2)})^2,$$

where $y^{(1)} = \frac{1}{|S_1|} \sum_{i \in S_1} y_i$, $y^{(2)} = \frac{1}{|S_2|} \sum_{i \in S_2} y_i$

Find the best split:

Try all the features & thresholds and find the one with **minimal objective function**

Parameters

- Maximum depth: (usually ~ 10)
- Minimum number of nodes in each node: (10, 50, 100)
- Single decision tree is not very powerful...
- Can we build multiple decision trees and ensemble them together?

Outline

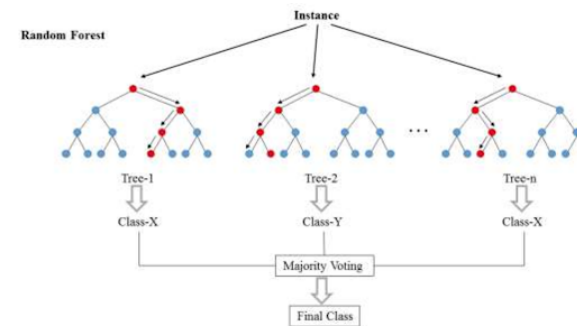
- Decision Tree
- **Random Forest**
- Gradient Boosted Decision Tree (GBDT)

Cho-Jui Hsieh, UC Davis ECS171: Machine Learning

Random Forest

- Random Forest (Bootstrap ensemble for decision trees):
 - Create T trees
 - Learn each tree using a subsampled dataset S_i and subsampled feature set D_i
 - Prediction: Average the results from all the T trees
- **Benefit:**
 - Avoid over-fitting
 - Improve stability and accuracy
- Good software available:
 - R: "randomForest" package Python: sklearn

An example



Building Decision Trees using MapReduce

- Parallel Learner for Assembling Numerous Ensemble Trees [Panda et al., VLDB '09]
 - A sequence of MapReduce jobs that builds a decision tree
 - Spark MLlib Decision Trees are based on PLANET

Outline

- Decision Tree
- Random Forest
- **Gradient Boosted Decision Tree (GBDT)**

Cho-Jui Hsieh, UC Davis ECS171: Machine Learning

Boosted Decision Tree

- Minimize loss $\ell(y, F(x))$ with $F(\cdot)$ being ensemble trees

$$F^* = \operatorname{argmin}_F \sum_{i=1}^n \ell(y_i, F(x_i)) \quad \text{with} \quad F(x) = \sum_{m=1}^T f_m(x)$$

(each f_m is a decision tree)

- Direct loss minimization: at each stage m , find the best function to minimize loss

- solve $f_m = \operatorname{argmin}_{f_m} \sum_{i=1}^N \ell(y_i, F_{m-1}(x_i) + f_m(x_i))$
- update $F_m \leftarrow F_{m-1} + f_m$

$F_m(x) = \sum_{j=1}^m f_j(x)$ is the prediction of x after m iterations.

- Two problems:
 - Hard to implement for general loss
 - Tend to overfit training data

Gradient Boosted Decision Tree (GBDT)

- Approximate the current loss function by a quadratic approximation:
用二次逼近法逼近电流损失函数:

$$\sum_{i=1}^n \ell_i(\hat{y}_i, f_m(x_i)) \approx \sum_{i=1}^n \left(\ell_i(\hat{y}_i) - g_i f_m(x_i) + \frac{1}{2} h_i f_m(x_i)^2 \right)$$

residual = actual - prediction

$$= \sum_{i=1}^n \frac{h_i}{2} \|f_m(x_i) - g_i/h_i\|^2 + \text{constant}$$

where $g_i = \partial_{\hat{y}_i} \ell_i(\hat{y}_i)$ is gradient,
 $h_i = \partial_{\hat{y}_i}^2 \ell_i(\hat{y}_i)$ is second order derivative

Gradient boosting (Freidman 1999)

Details

Gradient boosting (Friedman 1999) approximately solves (3) for arbitrary (differentiable) loss functions $\Psi(y, F(\mathbf{x}))$ with a two step procedure. First, the function $h(\mathbf{x}; \mathbf{a})$ is fit by least-squares

$$\mathbf{a}_m = \arg \min_{\mathbf{a}, \rho} \sum_{i=1}^N [\hat{y}_{im} - \rho h(\mathbf{x}_i; \mathbf{a})]^2 \quad (5)$$

to the current “pseudo”-residuals

$$\hat{y}_{im} = - \left[\frac{\partial \Psi(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(\mathbf{x})=F_{m-1}(\mathbf{x})} \quad (6)$$

Then, given $h(\mathbf{x}; \mathbf{a}_m)$, the optimal value of the coefficient β_m is determined

$$\beta_m = \arg \min_{\beta} \sum_{i=1}^N \Psi(y_i, F_{m-1}(\mathbf{x}_i) + \beta h(\mathbf{x}_i; \mathbf{a}_m)) \quad (7)$$

This strategy replaces a potentially difficult function optimization problem (3) by one based on least-squares (5), followed by a single parameter optimization (7) based on the general loss criterion Ψ .

Gradient Boosted Decision Tree (GBDT)

- Finding $f_m(\mathbf{x}, \theta_m)$ by minimizing the loss function:

$$\arg \min_{f_m} \sum_{i=1}^N [f_m(\mathbf{x}_i, \theta) - g_i/h_i]^2 + R(f_m)$$

- reduce the training of any loss function to regression tree (just need to compute g_i for different functions)
- $h_i = \alpha$ (fixed step size) for original GBDT.
- XGboost shows computing second order derivative yields better performance

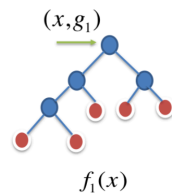
■ Algorithm:

Computing the current gradient for each \hat{y}_i .
Building a base learner (decision tree) to fit the gradient.
Updating current prediction $\hat{y}_i = F_m(\mathbf{x}_i)$ for all i .

Gradient Boosted Decision Tree (GBDT)

Key idea:

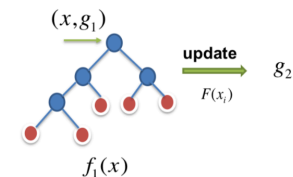
- Each base learner is a decision tree
- Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$



Gradient Boosted Decision Tree (GBDT)

Key idea:

- Each base learner is a decision tree
- Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$

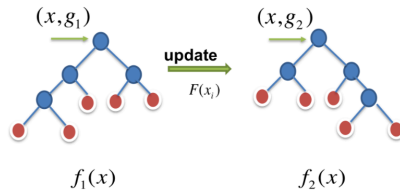


$$F_{m-1}(x_i) = \sum_{j=1}^{m-1} f_j(x_i) \quad g_m(x_i) = \left. \frac{\partial \ell(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x_i)=F_{m-1}(x_i)}$$

Gradient Boosted Decision Tree (GBDT)

Key idea:

- Each base learner is a decision tree
- Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$

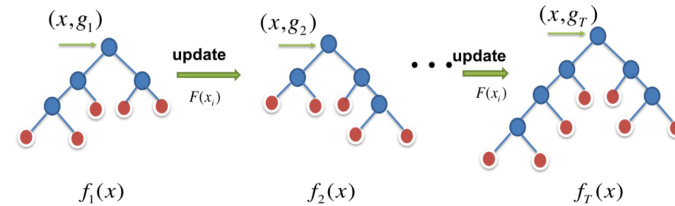


$$F_{m-1}(x_i) = \sum_{j=1}^{m-1} f_j(x_i) \quad g_m(x_i) = \left. \frac{\partial \ell(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x_i)=F_{m-1}(x_i)}$$

Gradient Boosted Decision Tree (GBDT)

Key idea:

- Each base learner is a decision tree
- Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$

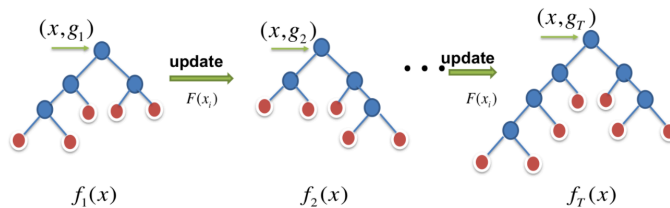


$$F_{m-1}(x_i) = \sum_{j=1}^{m-1} f_j(x_i) \quad g_m(x_i) = \left. \frac{\partial \ell(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x_i)=F_{m-1}(x_i)}$$

Gradient Boosted Decision Tree (GBDT)

Key idea:

- Each base learner is a decision tree
- Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$



Final prediction $F(x_i) = \sum_{j=1}^T f_j(x_i)$

Questions?